


Module IV

Process Management: Coordination And Synchronization

Location Of Process Coordination In The Hierarchy



Coordination Of Processes

- Is necessary in a concurrent system
- Avoids conflicts when multiple processes access shared items
- Allows a set of processes to cooperate
- Can also be used when
 - A process waits for I/O
 - A process waits for another process
- An example of cooperation among processes: UNIX pipes

Two Approaches To Process Coordination

- Use a hardware mechanism
 - Most useful /important on multiprocessor hardware
 - Often relies on *busy waiting*
- Use an operating system mechanism
 - Works well with single processor hardware
 - Does not entail unnecessary execution

Note: we will mention hardware quickly, and focus on operating system mechanisms

Two Key Situations That Process Coordination Mechanisms Handle

- Producer / consumer interaction
- Mutual exclusion

Producer-Consumer Synchronization

- Typical scenario: a FIFO buffer shared by multiple processes
 - Processes that deposit items into the buffer are called *producers*
 - Processes that extract items from the buffer are called *consumers*
- The programmer must guarantee
 - When the buffer is full, a producer must block until space is available
 - When the buffer is empty, a consumer must block until an item has been deposited
- A given process may act as a consumer for one buffer and a producer for another
- Example: a Unix pipeline

cat employees | grep Name: | sort

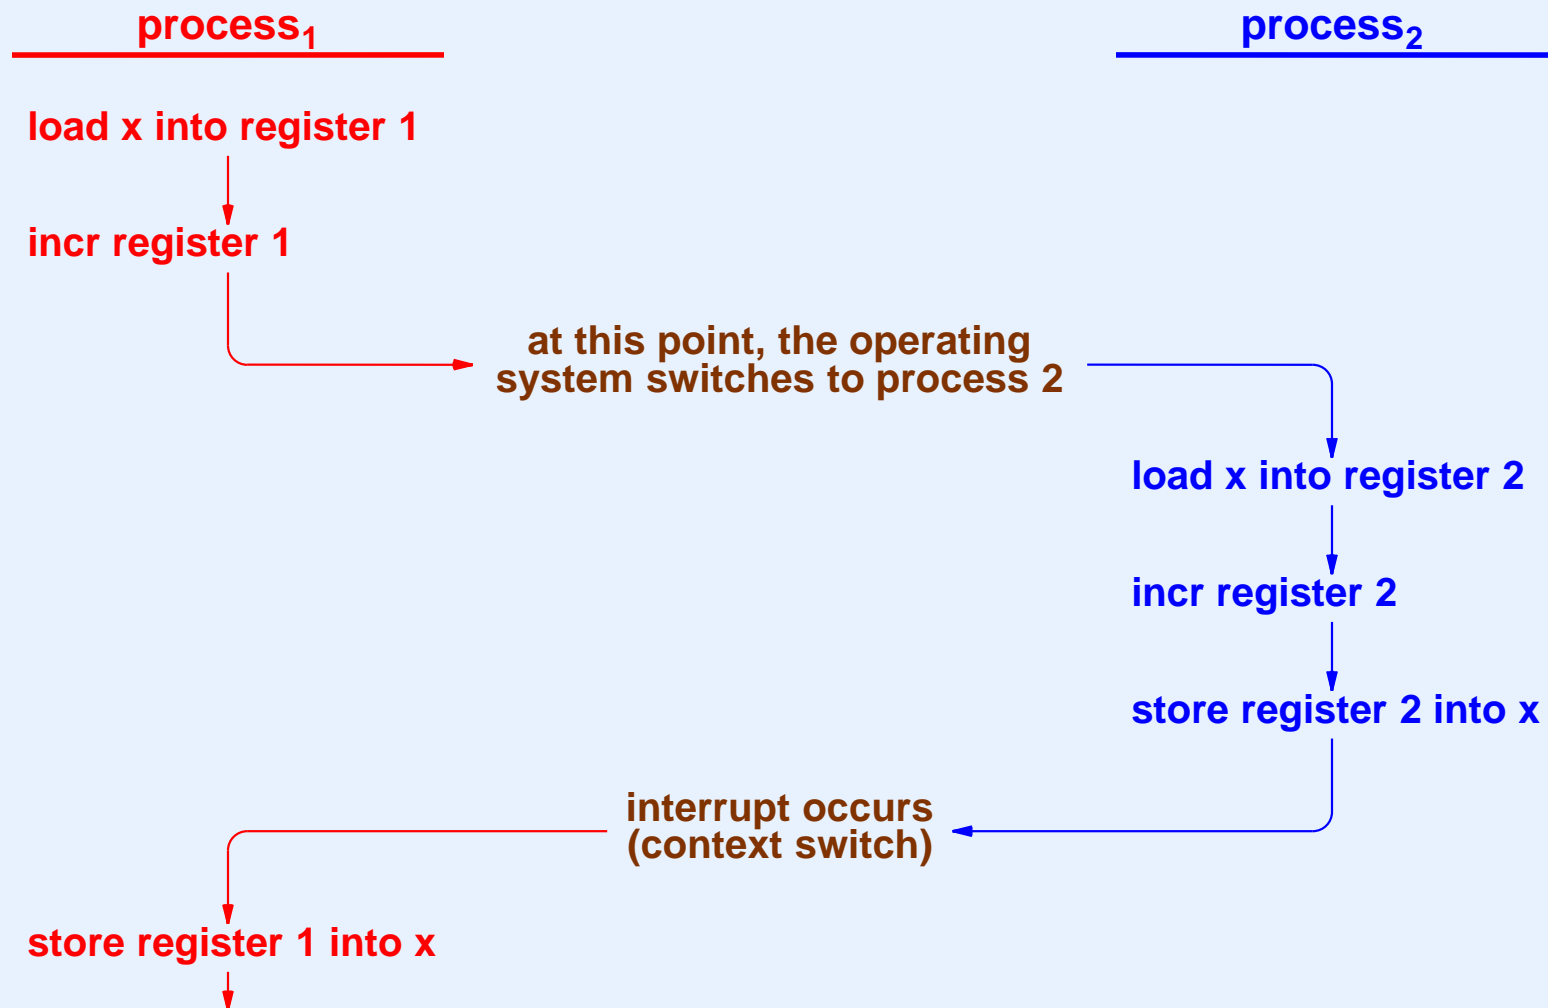
Mutual Exclusion

- In a concurrent system, multiple processes may attempt to access shared data items
- If one process starts to change a data item and then a context switch allows another process to run and access the data item, the results can be incorrect
- We use the term *atomic* to refer to an operation that is indivisible (i.e., the hardware performs the operation in a single instruction that cannot be interrupted)
- Programmers must take steps to ensure that when a sequence of non-atomic operations is used to change a data item, the sequence is not executed concurrently
- Even trivial operations can be non-atomic!

The Classic Example Of A Concurrent Access Problem

- Consider an integer, x
- To increment x , a programmer writes $x++$
- On most hardware architectures. three instructions are required
 - Load variable x into a register
 - Add 1 to the register
 - Store the register into variable x
- An operating system can switch from one process to another between any two instructions
- Surprising consequence: if two processes attempt to increment a shared integer concurrently, errors can result

Illustration Of What Can Happen When Two Processes Attempt To Increment Integer x Concurrently



To Prevent Problems

- A programmer must ensure that only one process accesses a shared item at any time
- General approach
 - Once a process obtains access, make all other processes wait
 - When a process finishes accessing the item, grant access to one of the waiting processes
- Three techniques are available
 - Hardware spin lock instructions
 - Hardware mechanisms that disable and restore interrupts
 - Semaphores (implemented in software)

Handling Mutual Exclusion With Spin Locks

- Used in multicore CPUs; does *not* work for a single processor
- An atomic hardware operation allows a core to test a memory location and change it
- The hardware guarantees that only one core will be allowed to make the change
- It is called a *spin lock* mechanism because a core uses *busy waiting* to gain access
- Busy waiting literally means the core executes a loop that tests the spin lock repeatedly until access is granted
- The approach is also known as *test-and-set*

An Example Of A Spin Lock (x86)

- An instruction performs an atomic compare and exchange (*cmpxchg*)
- Spin loop: repeat the following
 - Place an “unlocked” value (e.g., 0) in register *eax*
 - Place a “locked” value (e.g., 1) in register *ebx*
 - Place the address of a memory location in register *ecx* (the lock)
 - Execute the *cmpxchg* instruction
 - Register *eax* will contain the value of the lock before the compare and exchange occurred
 - Continue the spin loop as long as *eax* contains the “locked” value
- To release the lock, assign the “unlocked” value to the lock

Example Spin Lock Code For X86 (Part 1)

```
/* mutex.S - mutex_lock, mutex_unlock */

    .text
    .globl mutex_lock
    .globl mutex_unlock

/*-----
 * mutex_lock(uint32 *lock)  --  Acquire a lock
 *-----
 */
mutex_lock:

    /* Save registers that will be modified */

    pushl    %eax
    pushl    %ebx
    pushl    %ecx
```

Example Spin Lock Code For X86 (Part 2)

```
spinloop:
    movl    $0, %eax          /* Place the "unlocked" value in eax    */
    movl    $1, %ebx          /* Place the "locked" value in ebx     */
    movl    16(%esp), %ecx     /* Place the address of the lock in ecx */

    lock    cmpxchg %ebx, (%ecx) /* Atomic compare-and-exchange:        */
                                /* Compare %eax with memory (%ecx)     */
                                /* if equal                            */
                                /*      load %ebx in memory (%ecx)     */
                                /* else                                */
                                /*      load %ebx in %eax              */

    /* If eax is 1, the mutex was locked, so continue the spin loop */

    cmp     $1, %eax
    je      spinloop

    /* We hold the lock now, so pop the saved registers and return */
    popl    %ecx
    popl    %ebx
    popl    %eax
    ret
```

Example Spin Lock Code For X86 (Part 3)

```
/*-----  
 * mutex_unlock (uint32 *lock) - release a lock  
 *-----  
 */  
mutex_unlock:  
  
    /* Save register eax */  
    pushl    %eax  
  
    /* Load the address of lock onto eax */  
    movl     8(%esp), %eax  
  
    /* Store the "unlocked" value in the lock, thereby unlocking it */  
    movl     $0, (%eax)  
  
    /* Restore the saved register and return */  
    popl     %eax  
    ret
```

Handling Mutual Exclusion With Semaphores

- A programmer must allocate a semaphore for each item to be protected
- The semaphore acts as a *mutual exclusion* semaphore, and is known colloquially as a *mutex* semaphore
- All applications must be programmed to use the mutex semaphore before accessing the shared item
- The operating system guarantees that only one process can access the shared item at a given time
- The implementation avoids busy waiting

Definition Of Critical Section

- Each piece of shared data must be protected from concurrent access
- A programmer inserts mutex operations
 - Before access to the shared item
 - After access to the shared item
- The protected code is known as a *critical section*
- Mutex operations must be placed in each function that accesses the shared item

Mutual Exclusion Inside An Operating System

- Several possible approaches have been used
- Examples: allow only one process at a time to
 - Run operating system code
 - Run a given function
 - Access a given operating system component
- Allowing more processes to execute concurrently increases performance
- The general principle is: to maximize performance, choose the smallest possible granularity for mutual exclusion

Low-Level Mutual Exclusion

- Mutual exclusion is needed in two places
 - In application processes
 - Inside operating system
- On a single-processor system, mutual exclusion can be guaranteed provided that no context switching occurs
- A context switch can only occur when
 - A device interrupts
 - A process calls *resched*
- Low-level mutual exclusion technique: turn off interrupts and avoid rescheduling

Interrupt Mask

- A hardware mechanism that controls interrupts
- Implemented by an internal machine register, and may be part of *processor status word*
- On some hardware, a zero value means interrupts can occur; on other hardware, a non-zero value means interrupts can occur
- The OS can
 - Examine the current interrupt mask (find out whether interrupts are enabled)
 - Set the interrupt mask to prevent interrupts
 - Clear the interrupt mask to allow interrupts

Masking Interrupts

- Important principle:

No operating system function should contain code to explicitly enable interrupts.

- Technique used: a given function
 - Saves current interrupt status
 - Disables interrupts
 - Proceeds through a critical section
 - *Restores* the interrupt status from the saved copy
- Key insight: save / restore allows arbitrary call nesting

Why Interrupt Masking Is Insufficient

- It works! But...
- Stopping interrupts penalizes *all* processes when one process executes a critical section
 - It stops all I/O activity
 - It restricts execution to one process for the entire system
- Disabling interrupts can interfere with the scheduling invariant (e.g., a low-priority process can block a high-priority process for which I/O has completed)
- Disabling interrupts does not provide a policy that controls which process can access a critical section at a given time

High-Level Mutual Exclusion

- The idea is to create an operating system facility with the following properties
 - Permit applications to define multiple, independent critical sections
 - Allow processes to access each critical section independent of other sections
 - Provide an access policy that specifies how waiting processes gain access (e.g., FIFO)
- Good news: a single mechanism, the *counting semaphore*, suffices

Counting Semaphore

- An operating system abstraction
- An instance can be created dynamically
- Each instance is given a unique name
 - Typically an integer
 - Known as a *semaphore ID*
- An instance consists of a 2-tuple (count, set)
 - *Count* is an integer
 - *Set* is a set of processes that are waiting on the semaphore

Operations On Semaphores

- *Create* a new semaphore
- *Delete* an existing semaphore
- *Wait* on an existing semaphore
 - Decrements the count
 - Adds the calling process to set of waiting processes if the resulting count is negative
- *Signal* an existing semaphore
 - Increments the count
 - Makes a process ready if any are waiting

Xinu Semaphore Functions

`semid = semcreate(initial_count)` Creates a semaphore and returns an ID

`semdelete(semid)` Deletes the specified semaphore

`wait(semid)` Waits on the specified semaphore

`signal(semid)` signals the specified semaphore

Key Uses Of Counting Semaphores

- Semaphores have many potential uses
- However, using semaphores to solve complex coordination problems can be intellectually challenging
- We will consider two straightforward ways to use semaphores
 - Cooperative mutual exclusion
 - Producer-consumer synchronization (direct synchronization)

Cooperative Mutual Exclusion With Semaphores

- A set of processes use a semaphore to guard a shared item
- Initialize: create a mutex semaphore

```
sid = semcreate(1);
```

- Use: bracket each critical section in the code with calls to *wait* and *signal*

```
wait(sid);  
...critical section to use the shared item...  
signal(sid);
```

- All processes must agree to use semaphores (hence the term *cooperative*)
- Only one process will access the critical section at any time (others will be blocked)

A Potential Problem: Deadlock

- Consider two processes that use semaphores to protect two data items, x and y
- The two semaphores are created, and then the two processes run

```
sidx = semcreate(1);      sidy = semcreate(1);
```

- The processes take the following steps

```
/* Process 1 */
```

```
...  
wait(sidx);  
start to modify x  
wait(sidy);  
modify y  
signal(sidy);  
finish modifying x  
signal(sidx);
```

```
/* Process 2 */
```

```
...  
wait(sidy);  
start to modify y  
wait(sidx);  
modify x  
signal(sidx);  
finish modifying y  
signal(sidy);
```

Producer-Consumer Synchronization With Semaphores

- Two semaphores suffice to control processes accessing a shared buffer
- Initialize: create producer and consumer semaphores

```
psem = semcreate(buffer-size);  
csem = semcreate(0);
```

- The producer algorithm

```
repeat forever {  
    generate an item to be added to the buffer;  
    wait(psem);  
    fill_next_buffer_slot;  
    signal(csem);  
}
```

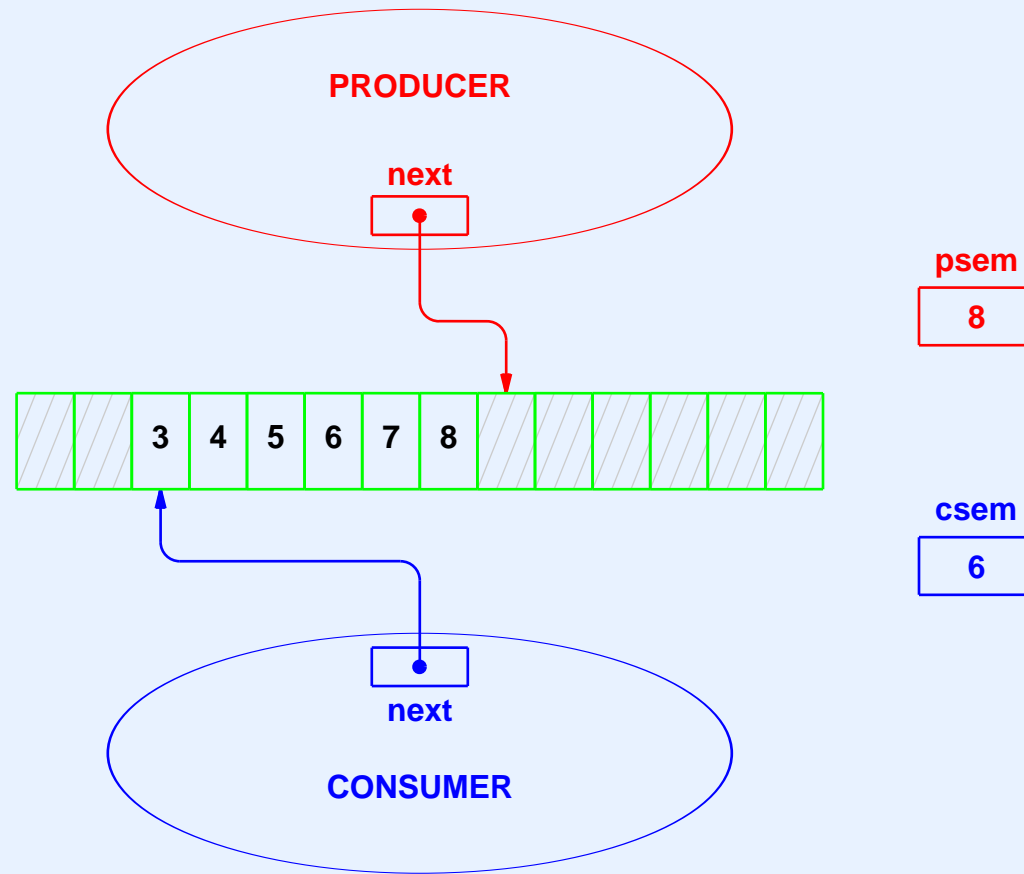
Producer-Consumer Synchronization With Semaphores

(continued)

- The consumer algorithm

```
repeat forever {  
    wait(csem);  
    extract_from_buffer_slot;  
    handle the item;  
    signal(psem);  
}
```

An Interpretation Of Producer-Consumer Semaphores



- *csem* counts the items currently in the buffer
- *psem* counts the unused slots in the buffer

The Semaphore Invariant

- Establishes a relationship between the semaphore concept and its implementation
- Makes the code easy to create and understand
- Must be re-established after each semaphore operation
- Is surprisingly elegant:

A nonnegative semaphore count means that the set of processes is empty. A count of negative N means that the set contains N waiting processes.


Counting Semaphores In Xinu

- Are stored in an array of semaphore entries
- Each entry
 - Corresponds to one instance (one semaphore)
 - Contains an integer count and pointer to a list of processes
- The ID of a semaphore is its index in the array
- The policy for management of waiting processes is FIFO

A Process State Used With Semaphores

- When a process is waiting on a semaphore, the process is not
 - Executing
 - Ready
 - Suspended
 - Receiving
- Note: the suspended state is only used by *suspend* and *resume*
- Therefore a new state is needed
- We will use the *WAITING* state for a process blocked by a semaphore

State Transitions With Waiting State



Semaphore Definitions

```
/* semaphore.h - isbadsem */

#ifndef NSEM
#define NSEM          120      /* Number of semaphores, if not defined */
#endif

/* Semaphore state definitions */

#define S_FREE    0      /* Semaphore table entry is available */
#define S_USED    1      /* Semaphore table entry is in use */

/* Semaphore table entry */
struct sentry {
    byte    sstate;      /* Whether entry is S_FREE or S_USED */
    int32    scount;      /* Count for the semaphore */
    qid16    squeue;      /* Queue of processes that are waiting */
                        /*      on the semaphore */
};

extern struct sentry semtab[];

#define isbadsem(s)      ((int32)(s) < 0 || (s) >= NSEM)
```

Implementation Of Wait (Part 1)

```
/* wait.c - wait */

#include <xinu.h>

/*-----
 * wait - Cause current process to wait on a semaphore
 *-----
 */
syscall wait(
    sid32      sem          /* Semaphore on which to wait */
)
{
    intmask mask;           /* Saved interrupt mask */
    struct procent *prpPtr; /* Ptr to process' table entry */
    struct sentry *semptr;  /* Ptr to semaphore table entry */

    mask = disable();
    if (isbadsem(sem)) {
        restore(mask);
        return SYSERR;
    }

    semptr = &semtab[sem];
    if (semptr->sstate == S_FREE) {
        restore(mask);
        return SYSERR;
    }
}
```

Implementation Of Wait (Part 2)

```
if (--(semptr->scount) < 0) {           /* If caller must block */
    prptr = &proctab[currpid];
    prptr->prstate = PR_WAIT;           /* Set state to waiting */
    prptr->prsem = sem;                 /* Record semaphore ID */
    enqueue(currpid,semptr->squeue);    /* Enqueue on semaphore */
    resched();                         /* and reschedule */
}

restore(mask);
return OK;
}
```

- Moving a process to the waiting state only requires a few lines of code
 - Set the state of the current process to PR_WAIT
 - Record the ID of the semaphore on which the process is waiting in field *prsem*
 - Call *resched*

The Semaphore Queuing Policy

- Determines which process to select among those that are waiting
- Is only used when *signal* is called and processes are waiting
- Examples of possible policies
 - First-Come-First-Served (FCFS or FIFO)
 - Process priority
 - Random

Semaphore Policy Consequences

- The goal is “fairness”
- Which semaphore queuing policy implements the goal the best?
- In other words, how should we interpret fairness?
- The semaphore policy can interact with scheduling policy
 - Should a low-priority process be allowed to access a resource if a high-priority process is also waiting?
 - Should a low-priority process be blocked forever if high-priority processes use a resource?

Choosing A Semaphore Queueing Policy

- The choice is difficult
- There is no single best answer
 - Fairness not easy to define
 - Scheduling and coordination interact in subtle ways
 - The choice may affect other OS policies
- The interactions of heuristic policies may produce unexpected results

The Semaphore Queuing Policy In Xinu

- First-come-first-serve
- Has several advantages
 - Is straightforward to implement
 - Is extremely efficient
 - Works well for traditional uses of semaphores
 - Guarantees all contending processes will obtain access
- Has an interesting consequence: a low-priority process can access a resource while a high-priority process remains blocked

Implementation Of A FIFO Semaphore Policy

- Each semaphore uses a list to manage waiting processes
- As we have seen Xinu manages the list of processes as a queue
 - Wait enqueues a process at one end of the queue
 - Signal chooses a process at the other end of the queue
- The code for signal follows

Implementation Of Signal (Part 1)

```
/* signal.c - signal */

#include <xinu.h>

/*-----
 * signal - Signal a semaphore, releasing a process if one is waiting
 *-----
 */
syscall signal(
    sid32      sem      /* ID of semaphore to signal */
)
{
    intmask mask;          /* Saved interrupt mask */
    struct sentry *semptr; /* Ptr to semaphore table entry */

    mask = disable();
    if (isbadsem(sem)) {
        restore(mask);
        return SYSERR;
    }
    semptr = &semtab[sem];
    if (semptr->sstate == S_FREE) {
        restore(mask);
        return SYSERR;
    }
}
```

Implementation Of Signal (Part 2)

```
    if ((semptr->scount++) < 0) {    /* Release a waiting process */  
        ready(dequeue(semptr->squeue));  
    }  
    restore(mask);  
    return OK;  
}
```

- Notice how little code is required to signal a semaphore

Possible Semaphore Allocation Strategies

- Static
 - All semaphores are defined at compile time
 - The approach is more efficient, but less powerful
- Dynamic
 - Semaphores are created at runtime
 - The approach is more flexible
- Xinu supports dynamic allocation, but preallocates the data structure to achieve efficiency

Xinu Semcreate (Part 1)

```
/* semcreate.c - semcreate, newsem */

#include <xinu.h>

local  sid32  newsem(void);

/*-----
 * semcreate - Create a new semaphore and return the ID to the caller
 *-----
 */
sid32  semcreate(
        int32          count          /* Initial semaphore count          */
)
{
    intmask mask;                    /* Saved interrupt mask          */
    sid32  sem;                      /* Semaphore ID to return        */

    mask = disable();

    if (count < 0 || ((sem=newsem())==SYSERR)) {
        restore(mask);
        return SYSERR;
    }
    semtab[sem].scount = count;      /* Initialize table entry        */

    restore(mask);
    return sem;
}
```


Xinu Semcreate (Part 2)

```
/*-----  
 * newsem - Allocate an unused semaphore and return its index  
 *-----  
 */  
local sid32 newsem(void)  
{  
    static sid32 nextsem = 0; /* Next semaphore index to try */  
    sid32 sem; /* Semaphore ID to return */  
    int32 i; /* Iterate through # entries */  
  
    for (i=0 ; i<NSEM ; i++) {  
        sem = nextsem++;  
        if (nextsem >= NSEM)  
            nextsem = 0;  
        if (semtab[sem].sstate == S_FREE) {  
            semtab[sem].sstate = S_USED;  
            return sem;  
        }  
    }  
    return SYSERR;  
}
```

Semaphore Deletion

- Wrinkle: one or more processes may be waiting when a semaphore is deleted
- We must choose how to dispose of each waiting process
- The Xinu disposition policy: if a process is waiting on a semaphore when the semaphore is deleted, the process becomes ready

Xinu Semdelete (Part 1)

```
/* semdelete.c - semdelete */

#include <xinu.h>

/*-----
 * semdelete - Delete a semaphore by releasing its table entry
 *-----
 */
syscall semdelete(
    sid32      sem      /* ID of semaphore to delete */
)
{
    intmask mask;        /* Saved interrupt mask */
    struct sentry *semptr; /* Ptr to semaphore table entry */

    mask = disable();
    if (isbadsem(sem)) {
        restore(mask);
        return SYSERR;
    }

    semptr = &semtab[sem];
    if (semptr->sstate == S_FREE) {
        restore(mask);
        return SYSERR;
    }
    semptr->sstate = S_FREE;
```

Xinu Semdelete (Part 2)

```
    resched_cntl(DEFER_START);  
    while (semptr->scount++ < 0) { /* Free all waiting processes */  
        ready(getfirst(semptr->squeue));  
    }  
    resched_cntl(DEFER_STOP);  
    restore(mask);  
    return OK;  
}
```

- Deferred rescheduling allows all waiting processes to be made ready before any of them to run
- Before it ends deferred rescheduling, semdelete ensures the semaphore data structure is ready for other processes to use



Do you understand semaphores?

Semaphore Behavior (A True Story)

- A process creates a semaphore

```
mutex = semcreate(1);
```

- Three processes then execute the following code

```
process convoy(char_to_print)
do forever {
    think (i.e., use CPU);
    wait(mutex);
    print(char_to_print);
    signal(mutex);
}
```

- The three processes print characters *A*, *B*, and *C*, respectively

The Convoy

- The initial output is
 - 20 *A*'s, 20 *B*'s, 20 *C*'s, 20 *A*'s, etc.
- After tens of seconds, however, the output becomes *ABCABCABC...*
- Facts
 - Everything is correct
 - No other processes are executing
 - The output is nonblocking (i.e., it uses polled I/O)

The Convoy

(continued)

- Questions
 - How long is thinking time?
 - Why does convoy start?
 - Will output switch back given enough time?
 - Did knowing the policies or the implementation of the scheduler and semaphore mechanisms make the convoy behavior obvious?

Summary

- Process synchronization is used in two ways
 - As a service supplied to applications
 - As an internal facility used inside the OS itself
- Low-level mutual exclusion
 - Masks hardware interrupts
 - Avoids rescheduling
 - Is insufficient for all coordination needs

Summary (continued)

- High-level process coordination is
 - Used by subsets of processes
 - Available inside and outside the OS
 - Implemented with counting semaphore
- Counting semaphore
 - A powerful abstraction implemented in software
 - Provides mutual exclusion and producer/consumer synchronization



Questions?