


Module V

Low-Level Memory Management Process Creation And Termination

Low-Level Memory Management

Location Of Low-Level Memory Management In The Hierarchy



The Apparent Impossibility Of A Hierarchical OS Design

- A process manager uses the memory manager to allocate space for a process
- A memory manager uses the device manager to page or swap to disk
- A device manager uses the process manager to block and restart processes when they request I/O
- Solution: divide the memory manager into two parts

The Two Types Of Memory Management

- Low-level memory manager
 - Manages memory within the kernel address space
 - Used to allocate address spaces for processes
 - Treats memory as a single, exhaustible resource
 - Positioned in the hierarchy below process manager
- High-level memory manager
 - Manages pages within a process's address space
 - Positioned in the hierarchy above the device manager
 - Divides memory into abstract resources

Conceptual Uses Of A Low-Level Memory Manager

- Allocate stack space for a process
 - Performed by the process manager when a process is created
 - The memory manager must include functions to allocate and free stacks
- Allocation of heap storage
 - Performed by the device manager (buffers) and other system facilities
 - The memory manager must include functions to allocate and free heap space

The Xinu Low-Level Memory Manager

- Two functions control allocation of stack storage

```
addr = getstk(numbytes);
```

```
freestk(addr, numbytes);
```

- Two functions control allocation of heap storage

```
addr = getmem(numbytes);
```

```
freemem(addr, numbytes);
```

- Memory is allocated until none remains
- Only *getmem/freemem* are intended for use by application processes; *getstk/freestk* are restricted to the OS

Well-Known Memory Allocation Strategies

- Stack and heap can be
 - Allocated from the same free area
 - Allocated from separate free areas
- The memory manager can use a single free list and follow a paradigm of
 - First-fit
 - Best-fit
 - The free list can be circular with a roving pointer
- The memory manager can maintain multiple free lists
 - By exact size (static / dynamic)
 - By range

Well-Known Memory Allocation Strategies

(continued)

- The free list can be kept in a hierarchical data structure (e.g., a tree)
 - Binary sizes of nodes can be used
 - Other sequences of sizes are also possible (e.g., Fibonacci)
- To handle repeated requests for the same size blocks, a cache can be combined with any of the above methods

Practical Considerations

- Sharing
 - A stack can never be shared
 - Multiple processes may share access to a given block allocated from the heap
- Persistence
 - A stack is associated with one process, and is freed when the process exists
 - An item allocated from a heap may persist longer than the process that created it
- Stacks tend to be one size, but heap requests vary in size
- Fragmentation can occur

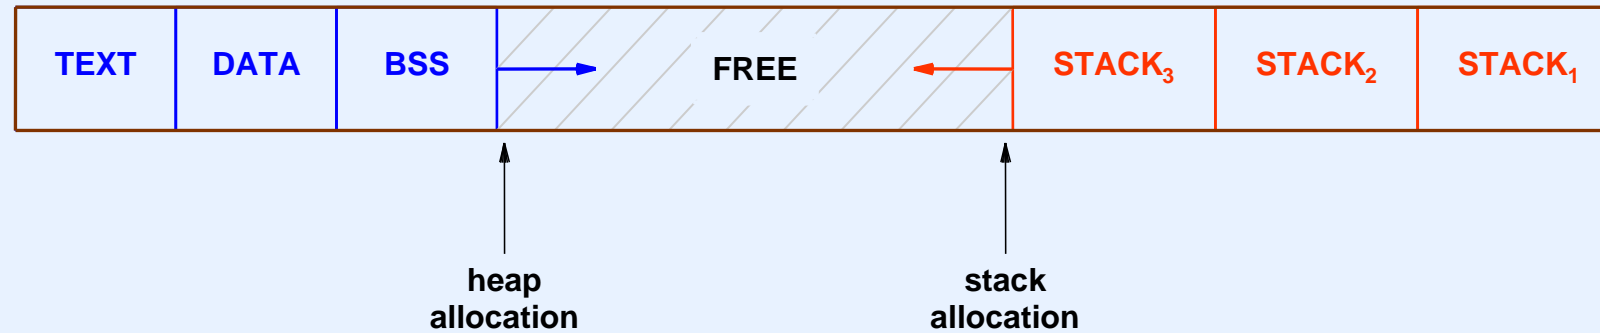
Memory Fragmentation

- Can occur if processes allocate and then free arbitrary-size blocks
- Symptom: after many requests to allocate and free blocks of memory, small blocks of allocated memory exist between blocks of free memory
- The problem: although much of the memory is free, each block on the free list is small
- Example
 - Assume a free memory consists of 1 Gigabyte total
 - A process allocates 1024 blocks of one Megabyte each (1 Gigabyte)
 - The process then frees every other block
 - Although 512 Megabytes of free memory are available, the largest free block is only 1 Megabyte

The Xinu Low-Level Allocation Scheme

- All free memory is treated as one resource
- A single free list is used for both heap and stack allocation
- The free list is
 - Ordered by increasing address
 - Singly-linked
 - Initialized at system startup to contain *all* free memory
- The Xinu allocation policies
 - Heap allocation uses the first-fit approach
 - Stack allocation uses the last-fit approach
 - The design results in two conceptual pools of memory

Consequence Of The Xinu Allocation Policy



- The first-fit policy means heap storage is allocated from lowest part of free memory
- The last-fit policy means stack storage is allocated from the highest part of free memory
- Note: because stacks tend to be uniform size, there is higher probability of reuse and lower probability of fragmentation

Protecting Against Stack Overflow

- Note that the stack for a process can grow downward into the stack for another
- Some memory management hardware supports protection
 - The memory for a process stack is assigned the process's protection key
 - When a context switch occurs the processor protection key is set
 - If a process overflows its stack, hardware will raise an exception
- If no hardware protection is available
 - Mark the top of each stack with a reserved value
 - Check the value when scheduling
 - The approach provides a little protection against overflow

Memory Allocation Granularity

- Facts
 - Memory is byte addressable
 - Some hardware requires alignment
 - * For process stack
 - * For I/ O buffers
 - * For pointers
 - Free memory blocks are kept on free list
 - One cannot allocate / free individual bytes
- Solution: choose a minimum granularity and round all requests to the minimum

Example Code To Round Memory Requests

```
/* excerpt from memory.h */

/*-----
 * roundmb, truncmb - Round or truncate address to memory block size
 *-----
 */
#define roundmb(x)      (char *) ( (7 + (uint32)(x)) & (~7) )
#define truncmb(x)      (char *) ( ((uint32)(x)) & (~7) )

struct memblk {
    struct memblk *mnext;    /* See roundmb & truncmb */
    uint32 mlength;          /* Ptr to next free memory blk */
};                           /* Size of blk (includes memblk) */

extern struct memblk memlist; /* Head of free memory list */
extern void *minheap;         /* Start of heap */
extern void *maxheap;         /* Highest valid heap address */
```

- Note the efficient implementation
 - The size of *memblk* is chosen to be a power of 2
 - The code implements rounding and truncation with bit manipulation

The Xinu Free List

- Employs a well-known trick: to link together a list of free blocks, place all pointers *in the blocks themselves*
- Each block on the list contains
 - A pointer to the next block
 - An integer giving the size of the block
- A fixed location (*memlist*) contains a pointer to the first block on the list
- Look again at the definitions in memory.h

Declarations For The Free List

```
/* excerpt from memory.h */


/*-----
 * roundmb, truncmb - Round or truncate address to memory block size
 *-----
 */
#define roundmb(x)      (char *) ( (7 + (uint32)(x)) & (~7) )
#define truncmb(x)      (char *) ( ((uint32)(x)) & (~7) )

struct memblk {          /* See roundmb & truncmb          */
    struct memblk *mnext; /* Ptr to next free memory blk */
    uint32 mlength;       /* Size of blk (includes memblk) */
};

extern struct memblk memlist; /* Head of free memory list */
extern void *minheap;        /* Start of heap */
extern void *maxheap;        /* Highest valid heap address */
```

- Struct *memblk* defines the two items stored in every block
- Variable *memlist* is the head of the free list
- Making the head of the list have the same structure as other nodes reduces special cases in the code

Illustration Of Xinu Free List



- Free memory blocks are used to store list pointers
- Items on the list are ordered by increasing address
- All allocations rounded to size of struct *memblk*
- The length in *memlist* counts total free memory bytes

Allocation Technique

- Round up the request to a multiple of memory blocks
- Walk the free memory list
- Choose either
 - First free block that is large enough (*getmem*)
 - Last free block that is large enough (*getstk*)
- If a free block is larger than the request, extract a piece for the request and leave the part that is left over on the free list

When Searching The Free List

- Use two pointers that point to two successive nodes on the list
- An invariant is used during the search
 - Pointer *curr* points to a node on the free list (or *NULL*)
 - Pointer *prev* points to the previous node (or *memlist*)
- The invariant is established initially by making *prev* point to *memblk* and making *curr* point to the item to which *memblk* points
- The invariant must be maintained each time pointers move along the list

Xinu Getmem (Part 1)

```
/* getmem.c - getmem */

#include <xinu.h>

/*-----
 * getmem - Allocate heap storage, returning lowest word address
 *-----
 */
char *getmem(
    uint32 nbytes /* Size of memory requested */
)
{
    intmask mask; /* Saved interrupt mask */
    struct memblk *prev, *curr, *leftover;

    mask = disable();
    if (nbytes == 0) {
        restore(mask);
        return (char *)SYSERR;
    }

    nbytes = (uint32) roundmb(nbytes); /* Use memblk multiples */
}
```

Xinu Getmem (Part 2)

```
prev = &memlist;
curr = memlist.mnext;
while (curr != NULL) {                                /* Search free list */

    if (curr->mlength == nbytes) { /* Block is exact match */
        prev->mnext = curr->mnext;
        memlist.mlength -= nbytes;
        restore(mask);
        return (char *)(curr);

    } else if (curr->mlength > nbytes) { /* Split big block */
        leftover = (struct memblk *)((uint32) curr +
                                       nbytes);
        prev->mnext = leftover;
        leftover->mnext = curr->mnext;
        leftover->mlength = curr->mlength - nbytes;
        memlist.mlength -= nbytes;
        restore(mask);
        return (char *)(curr);

    } else {                                           /* Move to next block */
        prev = curr;
        curr = curr->mnext;
    }
}
restore(mask);
return (char *)SYSERR;
}
```

Splitting A Block

- Occurs when *getmem* chooses a block that is larger than the requested size
- *Getmem* performs three steps
 - Compute the address of the piece that will be left over (i.e., the right-hand side of the block)
 - Link the leftover piece into the free list
 - Return the original block to the caller
- Note: the address of the leftover piece is $\text{curr} + \text{nbytes}$ (the addition must be performed using unsigned arithmetic because the high-order bit may be on)

Deallocation Technique

- Round up the specified size to a multiple of memory blocks (allows the user to specify the same value during deallocation that was used during allocation)
- Walk the free list, using *next* to point to a block on the free list, and *prev* to point to the previous block (or *memlist*)
- Stop when the address of the block being freed lies between *prev* and *next*
- Either: insert the block into the list or handle coalescing

Coalescing Blocks

- The term *coalescing* refers to the opposite of splitting
- Coalescing occurs when a block being freed is adjacent to an existing free block
- Technique: instead of adding the new block to the list, combine the new and existing block into one larger block
- Note: the code must check for coalescing with the preceding block, the following block, or both

Xinu Freemem (Part 1)

```
/* freemem.c - freemem */

#include <xinu.h>

/*-----
 * freemem - Free a memory block, returning the block to the free list
 *-----
 */
syscall freemem(
    char      *blkaddr,      /* Pointer to memory block */
    uint32     nbytes        /* Size of block in bytes */
)
{
    intmask mask;            /* Saved interrupt mask */
    struct memblk *next, *prev, *block;
    uint32 top;

    mask = disable();
    if ((nbytes == 0) || ((uint32) blkaddr < (uint32) minheap)
        || ((uint32) blkaddr > (uint32) maxheap)) {
        restore(mask);
        return SYSERR;
    }

    nbytes = (uint32) roundmb(nbytes); /* Use memblk multiples */
    block = (struct memblk *)blkaddr;
```

Xinu Freemem (Part 2)

```
prev = &memlist;                                /* Walk along free list */
next = memlist.mnext;
while ((next != NULL) && (next < block)) {
    prev = next;
    next = next->mnext;
}

if (prev == &memlist) {                          /* Compute top of previous block*/
    top = (uint32) NULL;
} else {
    top = (uint32) prev + prev->mlength;
}

/* Ensure new block does not overlap previous or next blocks */

if (((prev != &memlist) && (uint32) block < top)
    || ((next != NULL) && (uint32) block+nbytes>(uint32)next)) {
    restore(mask);
    return SYSERR;
}

memlist.mlength += nbytes;
```

Xinu Freemem (Part 3)

```
/* Either coalesce with previous block or add to free list */

if (top == (uint32) block) {      /* Coalesce with previous block */
    prev->mlength += nbytes;
    block = prev;
} else {                          /* Link into list as new node */
    block->mnext = next;
    block->mlength = nbytes;
    prev->mnext = block;
}

/* Coalesce with next block if adjacent */

if (((uint32) block + block->mlength) == (uint32) next) {
    block->mlength += next->mlength;
    block->mnext = next->mnext;
}
restore(mask);
return OK;
}
```

Xinu Getstk (Part 1)

```
/* getstk.c - getstk */

#include <xinu.h>

/*-----
 * getstk - Allocate stack memory, returning highest word address
 *-----
 */
char *getstk(
    uint32 nbytes /* Size of memory requested */
)
{
    intmask mask; /* Saved interrupt mask */
    struct memblk *prev, *curr; /* Walk through memory list */
    struct memblk *fits, *fitsprev; /* Record block that fits */

    mask = disable();
    if (nbytes == 0) {
        restore(mask);
        return (char *)SYSERR;
    }

    nbytes = (uint32) roundmb(nbytes); /* Use mblock multiples */

    prev = &memlist;
    curr = memlist.mnext;
    fits = NULL;
```

Xinu Getstk (Part 2)

```
while (curr != NULL) {                                /* Scan entire list */
    if (curr->mlength >= nbytes) {                    /* Record block address */
        fits = curr;                                  /* when request fits */
        fitsprev = prev;
    }
    prev = curr;
    curr = curr->mnext;
}

if (fits == NULL) {                                    /* No block was found */
    restore(mask);
    return (char *)SYSERR;
}

if (nbytes == fits->mlength) {                          /* Block is exact match */
    fitsprev->mnext = fits->mnext;
} else {                                                /* Remove top section */
    fits->mlength -= nbytes;
    fits = (struct memblk *)((uint32)fits + fits->mlength);
}

memlist.mlength -= nbytes;
restore(mask);
return (char *)((uint32) fits + nbytes - sizeof(uint32));
}
```

Xinu Freestk

```
/* excerpt from memory.h */

/*-----
 * freestk -- Free stack memory allocated by getstk
 *-----
 */
#define freestk(p,len) freemem((char *)((uint32)(p) \
                                - ((uint32)roundmb(len)) \
                                + (uint32)sizeof(uint32)), \
                                (uint32)roundmb(len) )
```

- Implemented as an inline function
- Technique: convert address from the highest address in block being freed to the lowest address in the block, and call *freemem*

Process Creation And Termination

Process Creation

- Process creation and termination use the memory manager
- Creation
 - Allocates a stack for the process being created
 - Fills in process table entry
 - Fills in the process's stack to have a valid frame
- Two design decisions
 - Choose an initial state for the process
 - Choose an action for the case where a process “returns” from the top-level function

The Xinu Design

- The initial state of a new process
 - A process is created in the suspended state
 - Consequence: execution can only begin after the process is resumed
- Return from top-level function
 - Causes the process to exit (similar to Unix)
 - Implementation: place a “pseudo call” on the stack (make it appear that the top-level function in the process was called)
 - Initialize the return address in the pseudo call to *INITRET*
- Note: *INITRET* is defined to be function *userret*
- Function *userret* causes the current process to exit

Xinu Function Userret

```
/* userret.c - userret */

#include <xinu.h>

/*-----
 * userret - Called when a process returns from the top-level function
 *-----
 */
void userret(void)
{
    kill(getpid());          /* Force process to exit */
}
```

The Pseudo Call On An Initial Stack

- Seems straightforward
- Is actually extremely tricky
- The trick: arrange the stack as if the new process was stopped in a call to *ctxsw*
- Several details make it difficult
 - *Ctxsw* runs with interrupts disabled, but a new process should start with interrupts enabled
 - We must store arguments for the new process so that the top-level function receives them
- We will examine code for process creation after looking at process termination

Killing A Process

- Formally known as *process termination*
- The action taken depends on the state of the process
 - If a process is on a list, it must be removed
 - If a process is waiting on a semaphore, the semaphore count must be adjusted
- In Xinu, function *kill* implements process termination

Xinu Implementation Of Kill (Part 1)

```
/* kill.c - kill */

#include <xinu.h>

/*-----
 * kill - Kill a process and remove it from the system
 *-----
 */
syscall kill(
    pid32      pid          /* ID of process to kill */
)
{
    intmask mask;           /* Saved interrupt mask */
    struct procent *prptr;  /* Ptr to process' table entry */
    int32 i;                /* Index into descriptors */

    mask = disable();
    if (isbadpid(pid) || (pid == NULLPROC)
        || ((prptr = &proctab[pid])->prstate) == PR_FREE) {
        restore(mask);
        return SYSERR;
    }

    if (--prcount <= 1) {   /* Last user process completes */
        xdone();
    }
}
```

Xinu Implementation Of Kill (Part 2)

```
send(prptr->prparent, pid);
freestk(prptr->prstkbase, prptr->prstklen);

switch (prptr->prstate) {
case PR_CURR:
    prptr->prstate = PR_FREE;          /* Suicide */
    resched();

case PR_SLEEP:
case PR_RECTIM:
    unsleep(pid);
    prptr->prstate = PR_FREE;
    break;

case PR_WAIT:
    semtab[prptr->prsem].scount++;
    /* Fall through */

case PR_READY:
    getitem(pid);                      /* Remove from queue */
    /* Fall through */

default:
    prptr->prstate = PR_FREE;
}

restore(mask);
return OK;
}
```


Killing The Current Process

- Look carefully at the code
 - Step 1: free the process's stack
 - Step 2: perform other actions
- Consider what happens when a current process kills itself: the call to *resched* occurs after the process's stack has been freed
- Why does it work?
- Answer: because in Xinu, even after stack has been freed, the memory is still available to the process

The Xdone Function

- Function *xdone* is called when the count of user processes reaches zero
- Nothing further will happen — only the null process remains running
- The function prints a warning message for the user

```
/* xdone.c - xdone */

#include <xinu.h>

/*-----
 * xdone - Print system completion message as last process exits
 *-----
 */
void xdone(void)
{
    kprintf("\n\nAll user processes have completed.\n\n");
    halt(); /* Halt the processor */
}
```

The Steps For Process Creation

- Allocate a process table entry
- Allocate a stack
- Place values on the stack as if the top-level function was called (pseudo-call)
- Arrange the saved state on the stack so context switch can switch to the process
- Details depend on
 - The hardware and calling conventions
 - The way context switch is written
- Consider example code for ARM and x86 processors

Process Creation On ARM (Part 1)

```
/* create.c - create, newpid */

#include <xinu.h>

local  int newpid();

#define roundew(x)      ( (x+3)& ~0x3)

/*-----
 * create - create a process to start running a procedure
 *-----
 */
pid32  create(
        void          *procaddr,      /* procedure address          */
        uint32        ssize,          /* stack size in bytes       */
        pri16         priority,       /* process priority > 0      */
        char          *name,          /* name (for debugging)      */
        uint32        nargs,          /* number of args that follow */
        ...
)
{
    intmask          mask;             /* interrupt mask             */
    pid32            pid;              /* stores new process id      */
    struct procent   *prptr;           /* pointer to proc. table entry */
    int32            i;
    uint32           *a;               /* points to list of args     */
    uint32           *saddr;           /* stack address              */

```

Process Creation On ARM (Part 2)

```
mask = disable();
if (ssize < MINSTK)
    ssize = MINSTK;
ssize = (uint32) roundew(ssize);
if (((saddr = (uint32 *)getstk(ssize)) ==
    (uint32 *)SYSERR) ||
    (pid=newpid()) == SYSERR || priority < 1 ) {
    restore(mask);
    return SYSERR;
}

prcount++;
prptr = &proctab[pid];

/* initialize process table entry for new process */
prptr->prstate = PR_SUSP;          /* initial state is suspended */
prptr->prprio = priority;
prptr->prstkbase = (char *)saddr;
prptr->prstklen = ssize;
prptr->prname[PNMLEN-1] = NULLCH;
for (i=0 ; i<PNMLEN-1 && (prptr->prname[i]=name[i])!=NULLCH; i++)
    ;
prptr->prsem = -1;
prptr->prparent = (pid32)getpid();
prptr->prhasmsg = FALSE;
```

Process Creation On ARM (Part 3)

```
/* set up initial device descriptors for the shell */
prpitr->prdesc[0] = CONSOLE; /* stdin is CONSOLE device */
prpitr->prdesc[1] = CONSOLE; /* stdout is CONSOLE device */
prpitr->prdesc[2] = CONSOLE; /* stderr is CONSOLE device */

/* Initialize stack as if the process was called */

*saddr = STACKMAGIC;

/* push arguments */
a = (uint32 *)(&nargs + 1); /* start of args */
a += nargs - 1; /* last argument */
for ( ; nargs > 4 ; nargs--) /* machine dependent; copy args */
    *--saddr = *a--; /* onto created process's stack */
*--saddr = (long)procaddr;
for(i = 11; i >= 4; i--)
    *--saddr = 0;
for(i = 4; i > 0; i--) {
    if(i <= nargs)
        *--saddr = *a--;
    else
        *--saddr = 0;
}
*--saddr = (long)INITRET; /* push on return address */
*--saddr = (long)0x00000053; /* CPSR F bit set, */
/* Supervisor mode */

prpitr->prstkptr = (char *)saddr;
restore(mask);
return pid;
}
```

Process Creation On ARM (Part 4)

```
/*-----  
 * newpid - Obtain a new (free) process ID  
 *-----  
 */  
local pid32 newpid(void)  
{  
    uint32 i; /* iterate through all processes */  
    static pid32 nextpid = 1; /* position in table to try or */  
                                /* one beyond end of table */  
  
    /* check all NPROC slots */  
  
    for (i = 0; i < NPROC; i++) {  
        nextpid %= NPROC; /* wrap around to beginning */  
        if (proctab[nextpid].prstate == PR_FREE) {  
            return nextpid++;  
        } else {  
            nextpid++;  
        }  
    }  
    return (pid32) SYSERR;  
}
```

Process Creation On X86 (Part 1)

```
/* create.c - create, newpid */

#include <xinu.h>

local  int newpid();

/*-----
 * create - Create a process to start running a function on x86
 *-----
 */
pid32  create(
        void          *funcaddr,      /* Address of the function      */
        uint32         ssize,          /* Stack size in bytes          */
        pri16          priority,       /* Process priority > 0         */
        char           *name,          /* Name (for debugging)         */
        uint32         nargs,          /* Number of args that follow   */
        ...
)
{
    uint32         savsp, *pushsp;
    intmask        mask;               /* Interrupt mask                */
    pid32          pid;                /* Stores new process id        */
    struct procent *prpPtr;            /* Pointer to proc. table entry */
    int32          i;
    uint32         *a;                 /* Points to list of args       */
    uint32         *saddr;             /* Stack address                 */
}
```


Process Creation On X86 (Part 2)

```
mask = disable();
if (ssize < MINSTK)
    ssize = MINSTK;
ssize = (uint32) roundmb(ssize);
if ( (priority < 1) || ((pid=newpid()) == SYSERR) ||
    ((saddr = (uint32 *)getstk(ssize)) == (uint32 *)SYSERR) ) {
    restore(mask);
    return SYSERR;
}

prcount++;
prptr = &proctab[pid];

/* Initialize process table entry for new process */
prptr->prstate = PR_SUSP;          /* Initial state is suspended */
prptr->prprio = priority;
prptr->prstkbase = (char *)saddr;
prptr->prstklen = ssize;
prptr->prname[PNMLEN-1] = NULLCH;
for (i=0 ; i<PNMLEN-1 && (prptr->prname[i]=name[i])!=NULLCH; i++)
    ;
prptr->prsem = -1;
prptr->prparent = (pid32)getpid();
prptr->prhasmsg = FALSE;
```

Process Creation On X86 (Part 3)

```
/* Set up stdin, stdout, and stderr descriptors for the shell */
prptr->prdesc[0] = CONSOLE;
prptr->prdesc[1] = CONSOLE;
prptr->prdesc[2] = CONSOLE;
/* Initialize stack as if the process was called */

*saddr = STACKMAGIC;
savsp = (uint32)saddr;

/* Push arguments */
a = (uint32 *)(&nargs + 1);      /* Start of args */
a += nargs - 1;                 /* Last argument */
for ( ; nargs > 0 ; nargs--)    /* Machine dependent; copy args */
    *--saddr = *a--;             /* onto created process' stack */
*--saddr = (long)INITRET;        /* Push on return address */
```

Process Creation On X86 (Part 4)

```
/* The following entries on the stack must match what ctxsw */
/* expects a saved process state to contain: ret address, */
/* ebp, interrupt mask, flags, registers, and an old SP */
*--saddr = (long)funcaddr; /* Make the stack look like it's */
/* half-way through a call to */
/* ctxsw that "returns" to the */
/* new process */
*--saddr = savsp; /* This will be register ebp */
/* for process exit */
savsp = (uint32) saddr; /* Start of frame for ctxsw */
*--saddr = 0x00000200; /* New process runs with */
/* interrupts enabled */

/* Basically, the following emulates an x86 "pushal" instruction */
*--saddr = 0; /* %eax */
*--saddr = 0; /* %ecx */
*--saddr = 0; /* %edx */
*--saddr = 0; /* %ebx */
*--saddr = 0; /* %esp; value filled in below */
pushsp = saddr; /* Remember this location */
*--saddr = savsp; /* %ebp (while finishing ctxsw) */
*--saddr = 0; /* %esi */
*--saddr = 0; /* %edi */
*pushsp = (unsigned long) (prptr->prstkptr = (char *)saddr);
restore(mask);
return pid;
}
```

Process Creation On X86 (Part 5)

```
/*-----  
 * newpid - Obtain a new (free) process ID  
 *-----  
 */  
local pid32 newpid(void)  
{  
    uint32 i; /* Iterate through all processes */  
    static pid32 nextpid = 1; /* Position in table to try or */  
                                /* one beyond end of table */  
  
    /* Check all NPROC slots */  
  
    for (i = 0; i < NPROC; i++) {  
        nextpid %= NPROC; /* Wrap around to beginning */  
        if (proctab[nextpid].prstate == PR_FREE) {  
            return nextpid++;  
        } else {  
            nextpid++;  
        }  
    }  
    return (pid32) SYSERR;  
}
```

An Assessment Of Process Creation

- Process creation code is among the most difficult code to understand
- One must know
 - The hardware architecture
 - The function calling conventions
 - The way *ctxsw* chooses to save state
 - How interrupts are handled
- As you struggle to understand it, imagine trying to write such code

Summary

- To preserve a multi-level hierarchy, the memory manager is divided into two pieces
 - A low-level manager is used in kernel to allocate address spaces
 - A high-level manager is used to handle abstractions of virtual memory and paging within a process's address space
- The Xinu low-level manager offers two types of allocation
 - Memory for a process stack
 - Memory from the heap
- Stack requests tend to repeat the same size

Summary (continued)

- The Xinu low-level memory manager
 - Places all free memory on a single list
 - Rounds all requests to multiples of *struct memblk*
 - Uses first-fit allocation for heap requests and last-fit allocation for stack requests
- Process creation and termination use the memory manager to allocate and free process stacks
- *Create* handcrafts an initial stack as if the top-level function had been called; the stack includes a return address given by constant *INITRET*



Questions?