


Module VII

High-Level Synchronous Message Passing

Location Of Synchronous Message Passing In The Hierarchy



A Review Of Xinu's Low-Level Message Passing Facility

- A message is always sent from one process directly to another
- Each process has a one-message message buffer
- Transmission is asynchronous (non-blocking)
- Reception is synchronous (blocking)
- An asynchronous function can be used to clear the message buffer

Features Of The Xinu High-Level Message Passing Mechanism

- Defines a set of message storage facilities called *ports*
- Specify the number of messages a given port can hold
- Many-to-many communication
 - Allow an arbitrary process to send a message to a port
 - Allow an arbitrary process to receive a message from a port
- Synchronous interface
 - Block a sender if a port is full
 - Block a receiver until a message arrives at a port
- Handle port deletion and reset

An Example Use Of Ports: A Concurrent Server

- Create a port, P
- Think of messages that are sent to the port as requests for some service
- Create a set of server processes that each repeatedly receive a request from P and “handle” the request (supply the service)
- Allow arbitrary processes to send requests to P
- Because server processes run concurrently, a server process can receive a later request and start handling it while another process continues to handle a previous request
- The advantage: short requests can be serviced quickly

A Few Details

- When the port system is initialized, a global pool of messages is created
 - The maximum number of messages in all ports is specified
 - Memory is allocated for the pool, and messages are linked onto a free list
- An individual port can be created (and later deleted) dynamically
- Semaphores are used to
 - Block a sender if the port is full
 - Block a receiver if the port is empty
- When a port is created
 - An argument specifies the number of messages that can be stored in the port
 - The request is denied if insufficient messages remain in the global pool

Functions That Operate On Ports

- *Ptinit*
 - Called once at startup
 - Initializes the entire port system
- *Ptcreate*
 - Creates a new port
 - An argument specifies maximum number of messages
- *Ptsend*
 - Sends a message to a port
- *Ptrecv*
 - Retrieves a message from a port

Functions That Operate On Ports

(continued)

- *Ptreset*
 - Resets existing port
 - Disposes of existing messages
 - Allows waiting processes to continue
- *Ptdelete*
 - Deletes existing port
 - Disposes of existing messages
 - Allows blocked processes to continue

Port Declarations

```
/* ports.h - isbadport */

#define NPORTS          30          /* Maximum number of ports */
#define PT_MSGS         100         /* Total messages in system */
#define PT_FREE         1           /* Port is free */
#define PT_LIMBO        2           /* Port is being deleted/reset */
#define PT_ALLOC        3           /* Port is allocated */

struct ptnode {                    /* Node on list of messages */
    uint32 ptmsg;                  /* A one-word message */
    struct ptnode *ptnext;         /* Pointer to next node on list */
};

struct ptenry {                    /* Entry in the port table */
    sid32 ptssem;                  /* Sender semaphore */
    sid32 ptrsem;                  /* Receiver semaphore */
    uint16 ptstate;                /* Port state (FREE/LIMBO/ALLOC) */
    uint16 ptmaxcnt;               /* Max messages to be queued */
    int32 ptseq;                   /* Sequence changed at creation */
    struct ptnode *pthread;         /* List of message pointers */
    struct ptnode *pttail;         /* Tail of message list */
};

extern struct ptnode *ptfree;      /* List of free nodes */
extern struct ptenry porttab[];    /* Port table */
extern int32 ptnextid;             /* Next port ID to try when
/*      looking for a free slot */

#define isbadport(portid)          ( (portid)<0 || (portid)>=NPORTS )
```

Xinu Ptnit (Part 1)

```
/* ptinit.c - ptinit */

#include <xinu.h>

struct ptnode *ptfree;          /* List of free message nodes */
struct ptnode porttab[NPORTS]; /* Port table */
int32 ptnextid;                 /* Next table entry to try */

/*-----
 * ptinit - Initialize all ports
 *-----
 */
syscall ptinit(
    int32 maxmsgs                /* Total messages in all ports */
)
{
    int32 i;                    /* Runs through the port table */
    struct ptnode *next, *curr; /* Used to build a free list */

    /* Allocate memory for all messages on all ports */

    ptfree = (struct ptnode *)getmem(maxmsgs*sizeof(struct ptnode));
    if (ptfree == (struct ptnode *)SYSERR) {
        panic("pinit - insufficient memory");
    }
}
```

Xinu Ptinit (Part 2)

```
/* Initialize all port table entries to free */

for (i=0 ; i<NPORTS ; i++) {
    porttab[i].ptstate = PT_FREE;
    porttab[i].ptseq = 0;
}
ptnextid = 0;

/* Create a free list of message nodes linked together */

for ( curr=next=ptfree ;  --maxmsgs > 0  ; curr=next ) {
    curr->ptnext = ++next;
}

/* Set the pointer in the final node to NULL */

curr->ptnext = NULL;
return OK;
}
```

Xinu Pcreate (Part 1)

```
/* pcreate.c - pcreate */

#include <xinu.h>

/*-----
 * pcreate - Create a port that allows "count" outstanding messages
 *-----
 */
syscall pcreate(
    int32      count      /* Size of port */
)
{
    intmask mask;          /* Saved interrupt mask */
    int32     i;           /* Counts all possible ports */
    int32     ptnum;       /* Candidate port number to try */
    struct    ptentry *ptptr; /* Pointer to port table entry */

    mask = disable();
    if (count < 0) {
        restore(mask);
        return SYSERR;
    }
}
```

Xinu Pcreate (Part 2)

```
for (i=0 ; i<NPORTS ; i++) {      /* Count all table entries      */
    ptnum = ptnextid;              /* Get an entry to check      */
    if (++ptnextid >= NPORTS) {
        ptnextid = 0;             /* Reset for next iteration  */
    }

    /* Check table entry that corresponds to ID ptnum */

    ptptr= &porttab[ptnum];
    if (ptptr->ptstate == PT_FREE) {
        ptptr->ptstate = PT_ALLOC;
        ptptr->ptssem = semcreate(count);
        ptptr->ptrsem = semcreate(0);
        ptptr->pthead = ptptr->pttail = NULL;
        ptptr->ptseq++;
        ptptr->ptmaxcnt = count;
        restore(mask);
        return ptnum;
    }
}
restore(mask);
return SYSERR;
}
```

Xinu Ptsend (Part 1)

```
/* ptsend.c - ptsend */

#include <xinu.h>

/*-----
 * ptsend - Send a message to a port by adding it to the queue
 *-----
 */
syscall ptsend(
    int32      portid,      /* ID of port to use          */
    umsg32     msg          /* Message to send           */
)
{
    intmask mask;           /* Saved interrupt mask      */
    struct ptentry *ptptr;  /* Pointer to table entry    */
    int32 seq;              /* Local copy of sequence num. */
    struct ptnode *msgnode; /* Allocated message node    */
    struct ptnode *tailnode; /* Last node in port or NULL */

    mask = disable();
    if ( isbadport(portid) ||
        (ptptr= &porttab[portid])->ptstate != PT_ALLOC ) {
        restore(mask);
        return SYSERR;
    }
}
```

Xinu Ptsend (Part 2)

```
/* Wait for space and verify port has not been reset */

seq = ptptr->ptseq;                /* Record original sequence */
if (wait(ptptr->ptssem) == SYSERR
    || ptptr->ptstate != PT_ALLOC
    || ptptr->ptseq != seq) {
    restore(mask);
    return SYSERR;
}
if (ptfree == NULL) {
    panic("Port system ran out of message nodes");
}

/* Obtain node from free list by unlinking */

msgnode = ptfree;                  /* Point to first free node */
ptfree = msgnode->ptnext;          /* Unlink from the free list */
msgnode->ptnext = NULL;            /* Set fields in the node */
msgnode->ptmsg = msg;
```

Xinu Ptsend (Part 3)

```
/* Link into queue for the specified port */

tailnode = ptptr->pttail;
if (tailnode == NULL) {          /* Queue for port was empty */
    ptptr->pttail = ptptr->pthead = msgnode;
} else {                         /* Insert new node at tail */
    tailnode->ptnext = msgnode;
    ptptr->pttail = msgnode;
}
signal(ptptr->ptrsem);
restore(mask);
return OK;
}
```


Xinu Precv (Part 1)

```
/* precv.c - precv */

#include <xinu.h>

/*-----
 * precv - Receive a message from a port, blocking if port empty
 *-----
 */
uint32 precv(
    int32 portid /* ID of port to use */
)
{
    intmask mask; /* Saved interrupt mask */
    struct ptentry *ptptr; /* Pointer to table entry */
    int32 seq; /* Local copy of sequence num. */
    umsg32 msg; /* Message to return */
    struct ptnode *msgnode; /* First node on message list */

    mask = disable();
    if ( isbadport(portid) ||
        (ptptr= &porttab[portid])->ptstate != PT_ALLOC ) {
        restore(mask);
        return (uint32)SYSERR;
    }
}
```

Xinu Precv (Part 2)

```
/* Wait for message and verify that the port is still allocated */

seq = ptptr->ptseq;                /* Record original sequence */
if (wait(ptptr->ptrsem) == SYSERR || ptptr->ptstate != PT_ALLOC
    || ptptr->ptseq != seq) {
    restore(mask);
    return (uint32)SYSERR;
}

/* Dequeue first message that is waiting in the port */

msgnode = ptptr->pthead;
msg = msgnode->ptmsg;
if (ptptr->pthead == ptptr->pttail) /* Delete last item */
    ptptr->pthead = ptptr->pttail = NULL;
else
    ptptr->pthead = msgnode->ptnext;
msgnode->ptnext = ptfree;           /* Return to free list */
ptfree = msgnode;
signal(ptptr->ptssem);
restore(mask);
return msg;
}
```

Port Deletion And Reset

- Illustrate how difficult it can be to delete resources in a concurrent system
- Situations that must be handled
 - If the port is full, processes may be blocked waiting to send messages to the port
 - If the port is empty, processes may be blocked waiting to receive messages from the port
 - If the port contains messages, some processing may be needed for each message
- An example of message processing during deletion
 - Suppose an application allocates heap memory and uses a message to send a pointer to the block of memory
 - When deleting such a port, the appropriate action may be to free the block of memory associated with each message

Disposing Of Messages

- Message disposition is needed during both reset and deletion
- What action should the system take to dispose of a message?
- Key idea: only the applications using the port will know how to dispose of messages
- To accommodate disposition
 - Both *ptreset* and *ptdelete* include an extra argument that specifies a disposition function
 - When a message is removed from the port, the disposition function is called with the message as an argument

How Dynamic Deletion Complicates A Design

- Consider what happens if
 - Process *A* invokes *ptsend* to send a message to a port
 - The port is full, so process *A* is blocked
 - While process *A* is blocked, process *B* starts to delete the port
 - Once the semaphores are deleted, process *A* will become ready
- If process *B* has lower priority than process *A*, process *A* will run
- How will process *A* know that the port is being deleted?
- A similar situation occurs for senders
- Another surprise: multiple processes may attempt to delete and/or reset the port concurrently

Concurrency And Message Disposition

- The function used to dispose of messages during deletion or reset
 - Is specified by user
 - May reschedule allowing other processes to execute
- An example
 - Suppose each message contains a pointer to a buffer from a buffer pool
 - The user's disposition function calls *freebuf* to free the buffer
 - *Freebuf* signals a semaphore, which calls *resched*
- Consequence: we need to handle attempts to use the port concurrently during reset or deletion

Three Mechanisms That Handle Port Reset

- Accession numbers: a sequence number is associated with each port
 - The sequence number is incremented when the port is created and when the port is deleted or reset
 - Functions *ptsend* and *ptrecv* record the sequence number when an operation begins and check the sequence number after *wait* returns
 - If sequence number changed, port was reset, so the operation should abort

Three Mechanisms For Handling Reset (continued)

- A new state for the port: assign each port a *state* variable
 - The value of the state variable indicates status
 - * *PTFREE* if the port is not in use
 - * *PTALLOC* if the port is in use
 - * *PTLIMBO* if the port is being reset/deleted
 - Functions *ptsend* and *ptrecv* examine the state variable
 - If the state is *PTLIMBO*, port is being reset or deleted and cannot be used

Three Mechanisms For Handling Reset (continued)

- Deferred rescheduling: temporarily postpone scheduling decisions
 - Call *resched_cntl(DEFER_START)* at start of reset or delete
 - Call *resched_cntl(DEFER_STOP)* after all operations are performed
- Note that deferred rescheduling means that message disposition will not start other concurrent processes

Xinu Ptdelete

```
/* ptdelete.c - ptdelete */

#include <xinu.h>

/*-----
 * ptdelete - Delete a port, freeing waiting processes and messages
 *-----
 */
syscall ptdelete(
    int32      portid,          /* ID of port to delete      */
    int32      (*disp)(int32)  /* Function to call to dispose */
    )              /* of waiting messages      */
{
    intmask mask;              /* Saved interrupt mask      */
    struct ptentry *ptptr;      /* Pointer to port table entry */

    mask = disable();
    if ( isbadport(portid) ||
        (ptptr= &porttab[portid])->ptstate != PT_ALLOC ) {
        restore(mask);
        return SYSERR;
    }
    _ptclear(ptptr, PT_FREE, disp);
    ptnextid = portid;
    restore(mask);
    return OK;
}
```

Xinu Ptreset

```
/* ptreset.c - ptreset */

#include <xinu.h>

/*-----
 * ptreset - Reset a port, freeing waiting processes and messages and
 *           leaving the port ready for further use
 *-----
 */
syscall ptreset(
    int32      portid,          /* ID of port to reset          */
    int32      (*disp)(int32)  /* Function to call to dispose */
    )              /* of waiting messages          */
{
    intmask mask;              /* Saved interrupt mask        */
    struct ptentry *ptptr;     /* Pointer to port table entry */

    mask = disable();
    if ( isbadport(portid) ||
        (ptptr= &porttab[portid])->ptstate != PT_ALLOC ) {
        restore(mask);
        return SYSERR;
    }
    _ptclear(ptptr, PT_ALLOC, disp);
    restore(mask);
    return OK;
}
```

Xinu _ptclear (Part 1)

```
/* ptclean.c - _ptclear */

#include <xinu.h>

/*-----
 * _ptclear - Used by ptdelete and ptreset to clear or reset a port
 *           (internal function assumes interrupts disabled and
 *           arguments have been checked for validity)
 *-----
 */
void _ptclear(
    struct ptentry *ptptr,          /* Table entry to clear */
    uint16 newstate,               /* New state for port */
    int32 (*dispose)(int32) /* Disposal function to call */
)
{
    struct ptnode *walk;           /* Pointer to walk message list */

    /* Place port in limbo state while waiting processes are freed */

    ptptr->ptstate = PT_LIMBO;

    ptptr->ptseq++;                 /* Reset accession number */
    walk = ptptr->pthead;           /* First item on msg list */
}
```

Xinu _ptclear (Part 2)

```
if ( walk != NULL ) {                                /* If message list nonempty */

    /* Walk message list and dispose of each message */

    for( ; walk!=NULL ; walk=walk->ptnext) {
        (*dispose)( walk->ptmsg );
    }

    /* Link entire message list into the free list */

    (ptptr->pttail)->ptnext = ptfree;
    ptfree = ptptr->pthead;
}

if (newstate == PT_ALLOC) {
    ptptr->pttail = ptptr->pthead = NULL;
    semreset(ptptr->ptssem, ptptr->ptmaxcnt);
    semreset(ptptr->ptrsem, 0);
} else {
    semdelete(ptptr->ptssem);
    semdelete(ptptr->ptrsem);
}
ptptr->ptstate = newstate;
return;
}
```

Summary

- Xinu offers a high-level message passing mechanism
- The system uses ports for message storage
- A port can be created dynamically, can have arbitrary senders, and arbitrary receivers
- The interface is completely synchronous — a sender blocks if a port is full, and a receiver blocks if a port is empty
- Port reset/deletion is tricky because
 - Concurrent processes may attempt to use the port while reset or deletion is occurring
 - Senders and receivers must be able to tell that the port changed while they were blocked

Summary (continued)

- Three techniques handle transition
 - A sequence number informs waiting processes whether the port was reset or deleted while they were blocked
 - A limbo state prevents new processes from using the port while it is being reset or deleted
 - Deferred rescheduling allows reset and deletion to proceed safely



Questions?