


Module XIV

User Interface

Location Of The User Interface In The Hierarchy



The Two Operating System Interfaces

- Operating systems provide two ways to access services
- An interface for applications
 - Generically called an API (Application Program Interface)
 - Consists of a set of *system calls*
- An interface for human users
 - Usually interactive
 - Can be a Command Line Interface (CLI) or Graphical User Interface (GUI)
 - Gives the system a “personality”

Characteristics Of User Interfaces

- GUI
 - Allows users to launch applications
 - May include *copy-and-paste* and *drag-and-drop* mechanisms
 - Relies on applications to handle most tasks
- CLI
 - Parses textual commands
 - Arguments on commands allow the user to specify an arbitrary level of detail

Command Interpreter Implementations

- Two implementations of command interpreters have been used
- The command interpreter is integrated into the operating system
- The command interpreter consists of an application that is separate from the operating system

Command Interpreter Built Into OS

- Typical of early systems and some small, embedded systems
- Advantages
 - The interpreter understands command syntax and semantics
 - Can offer command completion capability
 - Can prompt for arguments and check arguments for correctness
 - Can warn user about meaningless combinations
- Disadvantages
 - A user is limited to exactly the commands the OS provides
 - A user cannot select a non-standard command interpreter
 - Difficult to add commands (requires changing the OS)

Command Interpreter Implemented By An Application

- Introduced by MULTICS; popularized by Unix
- The interpreter only handles basic command syntax
- Individual commands must check and interpret arguments
- Advantage
 - Each user can choose their own interpreter
 - New commands can be added at any time
- Disadvantages
 - Nonuniformity among commands and arguments
 - No built-in semantic checks (argument errors are reported after a command starts running)

Example Of A Separate Interpreter: The Unix Shell

- Runs as a standard application (no special privilege is required)
- Provides per-line processing
- Interprets each line as a command
- Uses the same syntax for scripts as for interactive input
- Offers basic programming language constructs
 - Variables
 - Sequence of statements
 - Definite and indefinite iteration
 - Conditional execution

Shell Variables

- Have you used shell variables?
- Do you really understand how they work?
- The basics (from Korn shell)

```
X=hello  
echo $X
```

produces a line of output containing the word *hello*

- Given the above, the command

```
gcc $X.c
```

compiles file hello.c

Shell Binding Times

- Now consider a more complex example
- Suppose the current directory contains files

aaa bbb ccc

- What do the following lines mean if typed into a shell?

```
BEES=b*           # Define variable BEES
ls -l $BEES        # This will list file bbb
touch bb           # Add another file that starts with b
ls -l $BEES        # Will this line list just the file
                   #   named bbb, or both bb and bbb?
```

Another Example Of Binding Times

- The shell
 - Has both local and global variables
 - Uses the term *environment* for the set of global variables
- Environment variables
 - Imports variable definitions from the user's environment
 - Allows a user to *export* specific variables to the environment
 - The environment is passed to each child process that the shell executes
- Note: programs such as *make* allow environment variables to be accessed

Environment Variable Binding Times

- Suppose a user defines an environment variable `QQQ`

```
QQQ=CS354      # Define variable QQQ
export QQQ     # Export QQQ to the environment
myscript       # Run a shell script as a command
echo $QQQ      # Print the value of QQQ
```

- What will the output be if *myscript* contains the following lines?

```
echo $QQQ      # Print the current value of QQQ
QQQ=CS503      # Redefine QQQ
```

Environment Variable Binding Times (continued)

- The answer
 - A copy of the environment is kept for each process
 - A process inherits a copy from its parent when the process starts
 - Changes only affect the local copy (and processes that are created)
- In the example, the output is

CS354

CS354

The Basic Unix Shell Evaluation Algorithm

A shell repeats the following steps:

- A. Read and parse the next command, dividing it into tokens
- B. Perform variable assignment (*var = string*)
- C. Perform macro substitution: replace *\$X* with value of string *X*
- D. Perform file name matching (e.g., eliminate “*”)
- E. Search the user’s *PATH* for the command named by first token
- F. Invoke the command, passing remaining tokens as arguments

Data Types In The Unix Shell Language

- The shell supports one data type: string
- Builtin commands handle
 - Iteration (*while* and *for*)
 - Conditional execution (*if-then-else*)
- Quotes prevent substitution (delay binding)
 - Single quotes inhibit interpretation within the string
 - Double quotes allow variable substitution within the string
- Each command is executed by a separate process
- A command *pipeline* connects the output from one process to the input of another

Unix Shell Parsing In Practice

- The shell acts like a compiler
- Compound statements (*while*, *for*, *if-then-else* can span multiple lines of input
- A long pipeline can span multiple lines as well
- The shell must also handle file redirection
- Consequence: a shell must check for balanced delimiters (e.g., *if* → *fi*)

Unix Shell Data Conversions

- Output from a command can be converted to a string

``command``

- The contents of file can be assigned to a string

``cat file``

- The contents of a variable can be converted to command input

`echo $string | command`

- Literal text can be converted to command input

`command <<!`

...literal text goes here

!

The Unix Shell: Paths And Command Invocation

- The shell maintains “search path”
 - The path specifies a list of directories
 - The shell uses the path during command lookup
- To find a file to execute, the shell searches the path one directory at a time
 - It prepends the next directory to the command name
 - It checks to see if the result is a file
 - It stops if the file exists
- Once a file has been found, the shell checks to see that the file is executable
- The current directory (denoted “.”) works like any other directory name in a path

Efficiency And Path Binding

- There are two possibilities: late or early binding
- Late binding (used in the original shell)
 - The shell searches the path each time a user enters a command
 - Is inefficient
- Early binding (introduced in BSD Unix)
 - The shell presearches the path and caches the names of files in each directory
 - Is more efficient
 - Cannot detect changes in directory contents automatically
 - A user enters a *rehash* command to perform a new binding
 - The shell can automatically invoke *rehash* and retry if a command is not found

Efficiency And Path Binding

(continued)

- Does automatic rehash work?
- Answer: it depends
- Consider a case where
 - A user adds an executable program to a directory on the path
 - The program has a unique name (not found in any other directory along the path)
- Now consider a case where
 - A user adds an executable program to the first directory on the path
 - An executable program with the same name appears in a later directory on the path

I/O Redirection

- A user can redirect input or output
- The shell provides separate redirection for
 - Standard input
 - Standard output
 - Standard error
- Syntax is
 - Output: `> file`
 - Input: `< file`
- The syntax for standard error redirection depends on the shell

Synchronization Of Processes

- Background processing
 - The shell always creates a process to execute a command; background execution allows the shell to continue processing concurrently
 - The syntax is “&”
- Pipeline
 - The output from one process is fed to the input of another
 - An arbitrary pipeline is allowed
 - The “pipe” between two processes consists of a finite buffer
 - The operating system handles process synchronization

Shell Script

- The name given to a file that contains a set of shell commands
- The file must be executable
- A shell script uses the same syntax as an interactive shell (earlier operating systems used a special syntax for scripts)
- BSD Unix introduced the use of a two-byte *magic number* consisting of the ASCII characters `#!`
- If a file name follows the magic number, the file is taken to be the program to run with the script as input

`#! /users/me/bin/my_program`

Shell Input

- One can invoke a shell script, X , by:

`ksh < X`

or by naming X as an argument to the shell:

`ksh X`

- Challenge: create a shell script that behaves differently when invoked in the two ways shown above
- Note: feel free to use whatever shell you prefer (e.g., *bash*)

Design And Implementation Of An Example Shell

The Xinu Shell

- We will use the Xinu shell to illustrates the basics
- The Xinu shell
 - Is not fancy
 - Has a fixed set of commands compiled into shell itself
 - Uses a familiar command syntax:

`command_name arg* [redirection] [background]`

where *redirection* includes both input (< file) and output (> file), and *background* (&) allows a user to run a command in background

- The notation `arg*` means “zero or more args”, and `[x]` means “x is optional”
- The syntax and basic approach is similar to Unix

Lexical Tokens In The Xinu Shell

Token Type (num. value)	Character	Description
SH_TOK_AMPER (0)	&	ampersand
SH_TOK_LESS (1)	<	less-than symbol
SH_TOK_GREATER (2)	>	greater-than symbol
SH_TOK_OTHER (3)	'...'	quoted string (single quotes)
SH_TOK_OTHER (3)	"..."	quoted string (double-quotes)
SH_TOK_OTHER (3)	other	sequence of non-whitespace

- Only six lexical tokens are needed
- A string that starts with one type of quote can contain the other type of quote

"Don't blink!"

Organization Of Shell Software

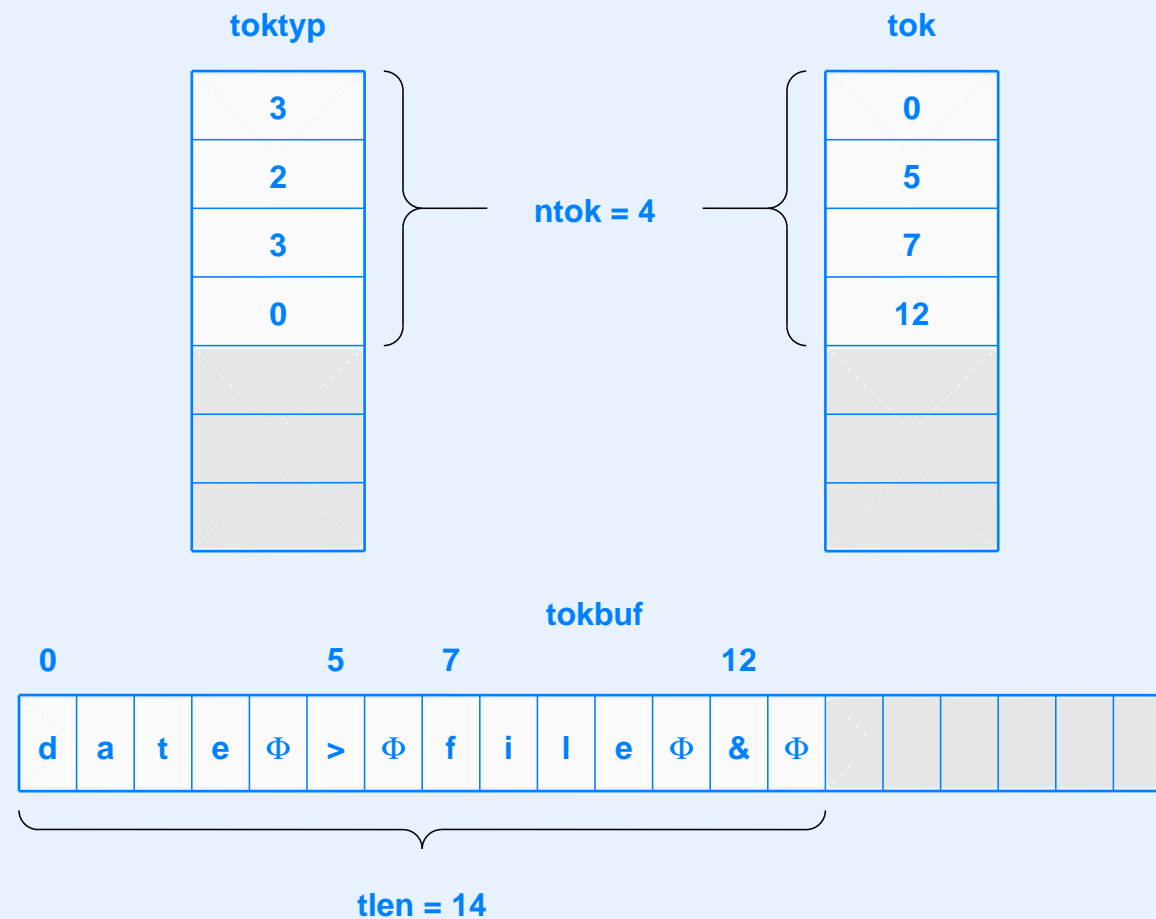
- The shell is organized like a compiler (interpreter)
- *A lexical analyzer*
 - Divides an input line into a series of tokens
 - Stores each token (the characters that make up the token) along with the type
- *A parser*
 - checks to ensure the sequence of tokens is valid
 - Turns the tokens into a command with arguments
 - executes the command

Lexical Analysis

- Because the Xinu shell only handles one line at a time, the shell
 - Reads an entire line
 - Calls a lexical analyzer to divide the line into tokens
- The lexical analyzer
 - Eliminates whitespace (i.e., blanks and tabs)
 - Checks for invalid tokens (e.g., file<)
 - Returns the number of tokens found

Token Storage

- Given the input line: `date > file &`
- The shell stores the tokens by storing an index and a type in two arrays




Parsing And Execution

- The first token must be the name of a command
- The final tokens may specify
 - Background processing
 - I/O redirection (input or output)
- The remaining tokens are taken to be arguments to the command

Arguments Passed To A Command

- Like Unix, Xinu only passes two arguments to a command
 - An integer count (nargs)
 - An array of pointers to argument strings
- Also like Unix, the first argument is the command name
- An example of arguments for input line `date -f illegal`



Builtin Commands

- Xinu allows a command to be *builtin*, which means the command is executed by the shell itself without creating a child process
- However
 - Builtin commands may not be run in background and may not have I/O redirection
- Examples
 - Exit (causes the shell to exit)
 - Kill (kill a process)
 - Clear (send the “screen clear” sequence to standard output)

Executing A Non-Builtin Command

- The shell must
 - Create a process to run the command
 - Construct an argument list and pass it to the command
 - Redirect input and/or output, if specified
 - Resume the process
- If the user did not specify background processing, the shell waits for the command to complete
- Note: when a process exits, Xinu sends a message to the parent with value equal to the process ID of the exiting process

A Few Shell Declarations (Part 1)

```
/* excerpt from shell.h */
```

```
/* Size constants */
```

```
#define SHELL_BUFLLEN    TY_IBUFLLEN+1    /* Length of input buffer      */
#define SHELL_MAXTOK     32               /* Maximum tokens per line    */
#define SHELL_CMDSTK     8192             /* Size of stack for process  */
                                           /*      that executes command */
#define SHELL_ARGLEN     (SHELL_BUFLLEN+SHELL_MAXTOK) /* Argument area              */
#define SHELL_CMDPRIO    20               /* Process priority for command */
```

```
/* Constants used for lexical analysis */
```

```
#define SH_NEWLINE      '\n'              /* New line character         */
#define SH_EOF          '\04'             /* Control-D is EOF          */
#define SH_AMPERS        '&'              /* Ampersand character        */
#define SH_BLANK         ' '              /* Blank character            */
#define SH_TAB          '\t'              /* Tab character              */
#define SH_SQUOTE        '\''             /* Single quote character     */
#define SH_DQUOTE        '"'              /* Double quote character     */
#define SH_LESS          '<'              /* Less-than character        */
#define SH_GREATER       '>'              /* Greater-than character     */
```

A Few Shell Declarations (Part 2)

```
/* Token types */

#define SH_TOK_AMPERSAND 0 /* Ampersand token */
#define SH_TOK_LESS 1 /* Less-than token */
#define SH_TOK_GREATER 2 /* Greater-than token */
#define SH_TOK_OTHER 3 /* Token other than those
                        /* listed above (e.g., an
                        /* alphanumeric string) */

/* Shell return constants */

#define SHELL_OK 0
#define SHELL_ERROR 1
#define SHELL_EXIT -3

/* Structure of an entry in the table of shell commands */

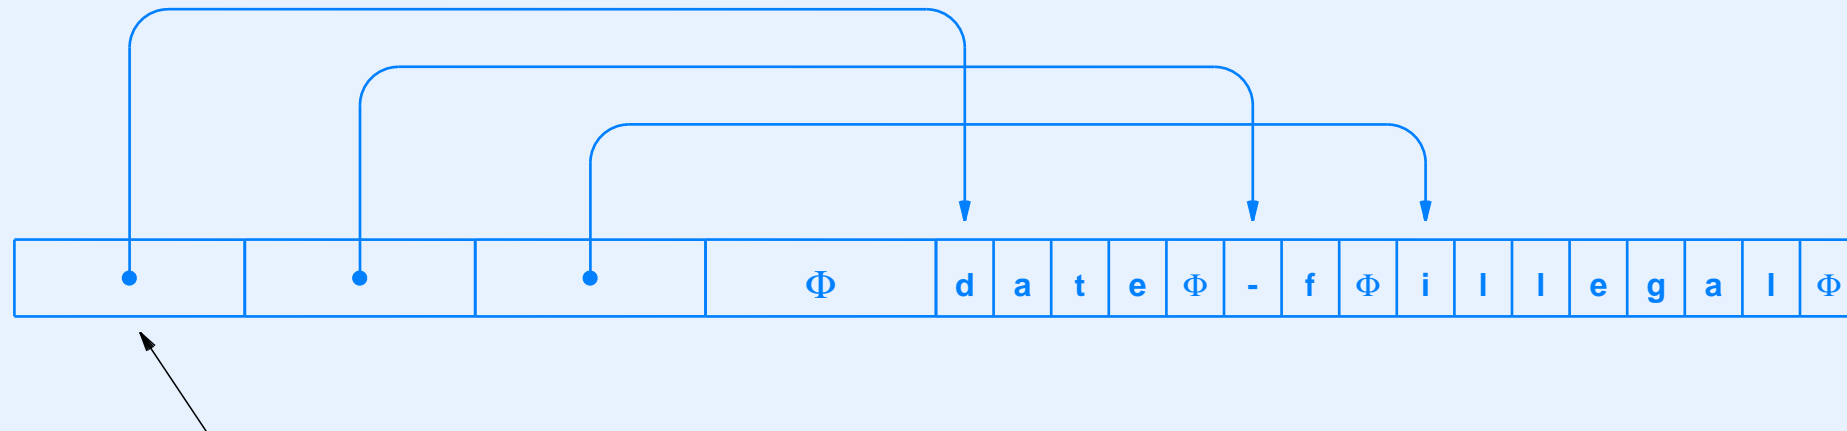
struct cmdent { /* Entry in command table */
    char *cname; /* Name of command */
    bool8 cbuiltin; /* Is this a builtin command? */
    int32 (*cfunc)(int32, char*[]); /* Function for command */
};

extern uint32 ncmd;
extern const struct cmdent cmdtab[];
```

Storage For Arguments

- Background processing means
 - Multiple commands can execute concurrently
 - Each command can have a different set of arguments
 - Arguments must be stored with the command process, not with the shell
- A Xinu trick: when creating a process to run a command, insert the arguments into bottom of the command's process stack
 - Store the *args* array plus the actual arguments
 - The arguments will be released along with the stack when the process exits
 - A single contiguous array is used
 - Added advantage: if a process accidentally overwrites its arguments, no other process will be affected

The Xinu Argument Layout At Runtime



args array starts at lowest byte of user's stack or next multiple of 4 bytes beyond it

- The shell
 - Creates a process to run a command, but does not resume the process
 - Computes the size needed for arguments, and copies them into the bottom of the process's stack
 - Resumes the process

For The Next Class

- Read Chapter 26 on the Xinu shell
- Look at
 - The lexical analysis function, *lexan.c*
 - The parser, *shell.c*
 - The function to handle arguments, *addargs.c*
- Come back with questions

Mice And Windows

A Mouse

- A pointing device invented as a companion to the bit-mapped display
- Operates as an I/O device
- The hardware
 - Detects movement
 - Reports motion in 2-dimensions
- The hardware interface is surprising

Mouse Hardware

- Contains
 - Two motion detectors
 - * Arranged at right angles
 - * Labeled X and Y
 - Two A-to-D converters
 - 1 to 3 buttons
 - Thumb wheels or touch controls
- Communicates with a computer using a character-oriented serial interface (the same interface as a keyboard)

Mouse Hardware

(continued)

- Mouse hardware is unexpectedly sophisticated
- It uses two-way communication to
 - Accept commands from the computer
 - Respond to queries from the computer
 - Transmit data to computer asynchronously

Mouse Communication And Interaction

- Two modes are used
 - Polling mode
 - Streaming mode

Polling Mode

- A processor sends a request for information
- The mouse responds by reporting
 - Motion since last request
 - The status of the buttons
- Seldom used except after communication has been temporarily lost

Streaming Mode

- A processor
 - Specifies the resolution and scaling to be used for motion detection
 - Sends one request to start the stream
- The mouse transmits new information
 - When movement exceeds a predetermined threshold
 - When a button is depressed or released
 - When thumbwheel moves

The Communications Interface

- A mouse uses a conventional serial line (or serial transmission over USB)
 - Full duplex
 - 8-bit characters
 - The RS232 standard can be used, just like a keyboard
- Each report sent by a mouse
 - Requires multiple characters
 - Follows a fixed format
 - The format is known as a *mouse packet*

Example Mouse Packet (2-Button Mouse)

YOVF	XOVF	YNEG	XNEG	reserved		BUT ₂	BUT ₁
reserved							
X motion since last report							
reserved							
Y motion since last report							
reserved							
Z motion (wheel) since last report							
reserved							

- Reserved bits must be zero
- *XOVF* and *YOVF* report overflow
- *XNEG* and *YNEG* are sign bits
- *BUT₁* and *BUT₂* give button status

Typical Mouse Parameters

- Samples per second
 - Selectable from 10, 20, 30,... 200 samples per second
 - Standard is 100 samples per second
- Tracking
 - Linear (e.g., 2:1)
 - Non-linear (e.g., exponential)
- Resolution
 - Example: 4 counts per millimeter
- The point: a seemingly simple device, a mouse, is quite complex

Display Screen

- A screen is divided into pixels
- Most screens are *memory mapped*, which means the operating system writes values to a display memory to change pixels on the screen
- Early screens were black-and-white
 - One bit in the memory per pixel
 - They were known as *bit-mapped* screens
- Current screens display color
 - One or more bytes per pixel
 - They are sometimes called *byte-mapped* screens

Color Specification

- The hardware uses three phosphors: red, green, blue
- All colors result from a combination of 0% to 100% of each of the three
 - The colors add in interesting ways
 - Example 100% red + 100% green + 0% blue gives yellow
- In practice
 - An integer value is used for each color instead of a percentage
 - Typically, the value for each color ranges from 0 to 255

Windowing Systems

- Windows are an operating system abstraction handled by software
- Most window systems allow windows to be
 - Created/destroyed at any time
 - Moved/resized/iconified
- A typical implementation
 - Each window is a rectangular region on screen
 - The window must be created before being used
 - The OS presents an application with a separate coordinate space where $(0,0)$ is the corner of the window

Window Display Parameters

- Many details are involved
 - Background/foreground colors
 - A title for the window
 - The location of scroll bars
 - Borders
- The details are controlled by a piece of software known as a *window manager*
- In Unix systems, each user can choose their own window manager
- Allowing users to choose a window manager means each user can see windows displayed in their preferred style

Cursor Movement

- The goal
 - A cursor should appear on the screen at all times
 - The cursor on the screen should track the mouse/touchpad movement
- Unfortunately
 - Mouse hardware is not directly linked to video hardware
- Consequence: software must
 - Update the cursor when the mouse moves
 - Map the new position to the correct window

The Cursor Update Algorithm

- When the cursor moves
 - Undo the cursor at old position by repainting the original display values
 - Determine the new position for the cursor, P
 - Save video memory at position P for a later “undo” operation
 - Use cursor color to paint cursor at position P

Optimizing Cursor Update

- A cursor update is required at each mouse interrupt
- Switching context to a window manager process introduces significant delay
- To reduce the overhead and avoid delay, perform the cursor update in the device driver during interrupt processing
- The downside: if the processor becomes overloaded, interrupt processing can be delayed or missed, which means the cursor on the screen may lag mouse movements or some mouse movements may be missed
- Good news: multicore processors make lost interrupts unlikely

Summary

- Operating system support two styles of user interface
 - Command line interpreter
 - Graphical
- The Unix shell
 - Runs as a separate application
 - Provides a miniature programming language
 - Supports concurrency and data pipelining
 - Limits variables to strings
 - Uses quotes to delay binding
 - Provides conversions between strings and command input/output

Summary (continued)

- A mouse
 - Is surprisingly complex
 - Uses a two-way communication system and delivers “packets”
- Display Screen
 - Uses a memory-mapped approach
 - The hardware for a pixel uses three primary colors: red, green, and blue

Summary (continued)

- Windowing systems
 - Are a software abstraction
 - Window manager handles details
- Cursor update
 - Is performed by software
 - Is often handled at interrupt time



Questions?